

TMA Training Center (TTC)

Java Code Optimization

<i>Course</i>	Java Programming
<i>Trainer</i>	Tuan Le
<i>Designed by</i>	TTC
<i>Last updated</i>	6-Nov-14

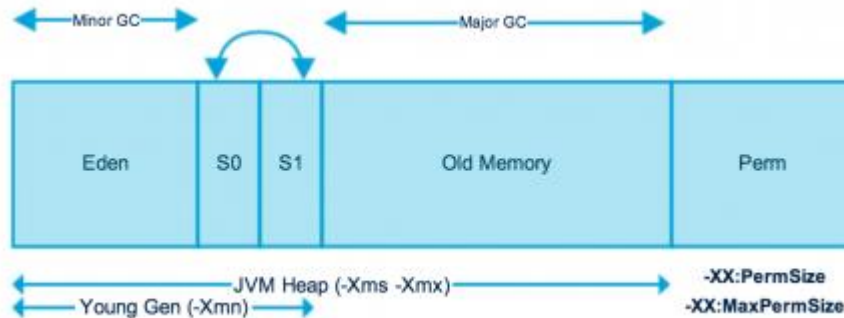
Contents

- Memory management & garbage collection
- Effective Java
- Java Optimization Tips

Objectives

- The course helps attendees to obtain some techniques that helps improve performance in java program.

Java (JVM) Memory Model



- JVM Heap memory is physically divided into two parts – **Young Generation** and **Old Generation**.

Java (JVM) Memory Model

- Young Generation Spaces:
 - Most of the newly created objects are located in the Eden memory space.
 - When Eden space is filled with objects, Minor GC is performed and all the survivor objects are moved to one of the survivor spaces.
 - Objects that are survived after many cycles of GC, are moved to the Old generation memory space.
- Old Generation
 - Contains the objects that are long lived and survived after many rounds of Minor GC.
 - Garbage collection is performed in Old Generation memory when it's full.
 - Old Generation Garbage Collection is called **Major GC** and usually takes longer time.
- Permanent Generation
 - The method area is implemented as a separated part.

Java memory management

- Objects are created on heap in Java irrespective of their scope e.g. local or member variable
- Class variables or static members are created in method area and both heap and method area is shared between different thread.
- Java lets you allocate objects as necessary and trust that they'll be reclaimed and recycled by the JVM
- The basic principle of garbage collection is the same in all cases:
 - Identify objects that are no longer in use by the program.
 - Recycle the memory used by these objects to create new ones.

Java memory management

- Define Java heap space for your application
 - In Eclipse IDE, if your program is consuming a lot of memory (loading big data) like this :

```
List<Domain> list = domainBo.findAllDomain(100000);  
for(Domain domain : list){  
    process(domain.getDomainName());  
}
```

- It can easily hit java.lang.OutOfMemoryError: Java heap space :
Exception in thread "main" java.lang.OutOfMemoryError: Java heap space
at java.util.HashMap.<init>(HashMap.java:209)
at java.util.LinkedHashMap.<init>(LinkedHashMap.java:181)

Java memory management

- Define Java heap space for your application

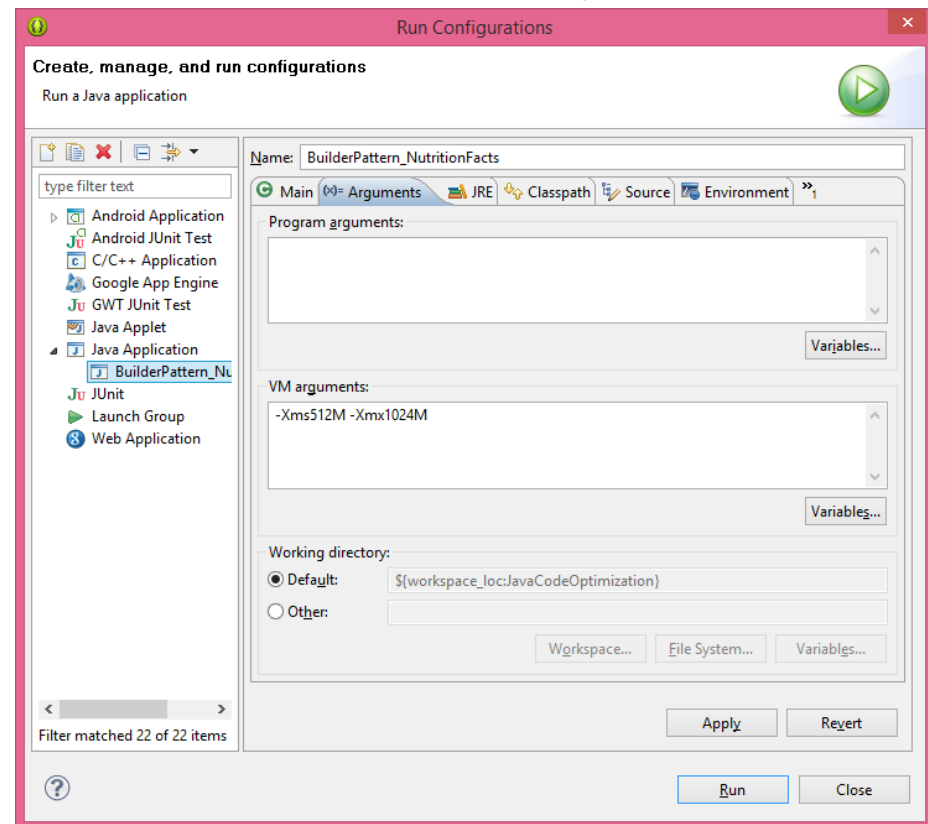
- Solution – VM arguments

On Eclipse menu, clicks Run -> Run Configurations..., select the Java application you want to run, clicks on the Arguments tab:

-Xms<size> - Set initial Java heap size

-Xmx<size> - Set maximum Java heap size

For example, -Xms512M -Xmx1024M



Java memory management

- Define Java heap space for your application using command line

- **Java Heap Size**

- ```
$ java -Xms512m -Xmx1024m JavaApp
```

- **Perm Gen Size**

- ```
$ java -XX:PermSize=64m -XX:MaxPermSize=128m JavaApp
```

- **Java Stack Size**

- ```
$ java -Xss512k JavaApp
```

# When an object becomes eligible for Garbage Collection

- Object becomes eligible for garbage collection if its *all references are null*.
- Cyclic dependencies are not counted as reference
  - If Object A has reference of object B and object B has reference of Object A and they don't have any other live reference
- Object is created inside a block and reference goes out scope once control exit that block.
- Parent object set to null.

# Static factory methods (1/3)

- A class can provide a public *static factory method* that returns an instance of the class

```
public static Boolean valueOf(boolean b) {
 return b ? Boolean.TRUE : Boolean.FALSE;
}
```

- Advantages:
  - Unlike constructors, they have names
  - They are not required to create a new object each time they're invoked
  - They can return an object of any subtype of their return type

# Static factory methods (2/3)

## ■ Advantages:

- They reduce the verbosity of creating parameterized type instances
- This typically requires you to provide the type parameters twice in quick succession:

```
Map<String, List<String>> m = new HashMap<String, List<String>>();
```

- With static factories, the compiler can figure out the type parameters for you

```
public static <K, V> HashMap<K, V> newInstance() {
 return new HashMap<K, V>();
}
```

- Then you could replace the wordy declaration above with this succinct alternative:

```
Map<String, List<String>> m = HashMap.newInstance();
```

# Static factory methods (3/3)

- Disadvantages:
  - Classes without public or protected constructors cannot be subclasses
  - They are not readily distinguishable from other static methods. Here are some common names for static factory methods:
    - `valueOf`, `getInstance`, `newInstance`, `getType`, `newType`

# Builder Pattern

- Static factories and constructors share a limitation: they do not scale well to large numbers of parameters: *telescoping constructor* pattern

```
class NutritionFacts {
 private final int servingSize; // (mL) required
 private final int servings; // (per container) required
 private final int calories; // optional
 private final int fat; // (g) optional
 private final int sodium; // (mg) optional
 private final int carbohydrate; // (g) optional
 public NutritionFacts(int servingSize, int servings) {
 this(servingSize, servings, 0);
 }
 public NutritionFacts(int servingSize, int servings, int calories) {
 this(servingSize, servings, calories, 0);
 }
 public NutritionFacts(int servingSize, int servings, int calories, int fat) {
 this(servingSize, servings, calories, fat, 0);
 }
 public NutritionFacts(int servingSize, int servings, int calories, int fat,
 int sodium) {
 this(servingSize, servings, calories, fat, sodium, 0);
 }
 public NutritionFacts(int servingSize, int servings, int calories, int fat,
 int sodium, int carbohydrate) {
 this.servingSize = servingSize;
 this.servings = servings;
 this.calories = calories;
 this.fat = fat;
 this.sodium = sodium;
 this.carbohydrate = carbohydrate;
 }
}
```

# Builder Pattern

- The telescoping constructor pattern works, but it is hard to write client: using JavaBeans pattern
  - JavaBean may be in an inconsistent state partway through its construction

```
class NutritionFacts {
 private int servingSize = -1; // Required; no default value
 private int servings = -1; // " " " "
 private int calories = 0;
 private int fat = 0;
 private int sodium = 0;
 private int carbohydrate = 0;
 public NutritionFacts() { }
 // Setters
 public void setServingSize(int val) { servingSize = val; }
 public void setServings(int val) { servings = val; }
 public void setCalories(int val) { calories = val; }
 public void setFat(int val) { fat = val; }
 public void setSodium(int val) { sodium = val; }
 public void setCarbohydrate(int val) { carbohydrate = val; }
}
```

# Builder Pattern

- Combines the safety of the telescoping constructor pattern with the readability of the JavaBeans pattern: *Builder* pattern

```
class NutritionFacts {
 private final int servingSize;
 private final int servings;
 private final int calories;
 private final int fat;
 private final int sodium;
 private final int carbohydrate;
 public static class Builder {
 // Required parameters
 private final int servingSize;
 private final int servings;
 // Optional parameters - initialized to default values
 private int calories = 0;
 private int fat = 0;
 private int carbohydrate = 0;
 private int sodium = 0;
 public Builder(int servingSize, int servings) {
 this.servingSize = servingSize;
 this.servings = servings;
 }
 }
}
```



# Builder Pattern

```
public Builder calories(int val) {calories = val;
 return this;
}
public Builder fat(int val) {
 fat = val;
 return this;
}
public Builder carbohydrate(int val) {
 carbohydrate = val;
 return this;
}
public Builder sodium(int val) {
 sodium = val;
 return this;
}
public NutritionFacts build() {
 return new NutritionFacts(this);
}
}

private NutritionFacts(Builder builder) {
 servingSize = builder.servingSize;
 servings = builder.servings;
 calories = builder.calories;
 fat = builder.fat;
 sodium = builder.sodium;
 carbohydrate = builder.carbohydrate;
}
}
```

# Builder Pattern

- This client code is easy to write and, more importantly, to read

```
NutritionFacts cocaCola = new NutritionFacts.Builder(240, 8).
calories(100).sodium(35).carbohydrate(27).build();
```

# Singleton pattern using enum type

- Making constructor private and having another factory method to get instance

```
public class OldSingleton {
 private static OldSingleton instance = null;
 private OldSingleton() {
 }
 public static OldSingleton getInstance() {
 if(instance == null)
 instance = new OldSingleton();
 return instance;
 }
}
```

# Singleton pattern using enum type

- Making constructor private and having another factory method to get instance:

```
public class OldSingleton {
 private static final OldSingleton instance = new OldSingleton();
 private OldSingleton() {
 }
 public static OldSingleton getInstance() {
 return instance;
 }
}
```

- First one is not thread safe and the second creates the Singleton Object even before it is actually required.

# Singleton pattern using enum type

- Refine our first approach to make it thread safe

```
public class OldSingleton {
 private static OldSingleton instance = null;
 private OldSingleton() {
 }
 public static synchronized OldSingleton getInstance(){
 if(instance == null)
 instance = new OldSingleton();
 return instance;
 }
}
```

# Singleton pattern using enum type

- Do one more refinement by fine grained locking

```
public class OldSingleton {
 private static OldSingleton instance = null;
 private OldSingleton() {
 }
 public static OldSingleton getInstance(){
 if(instance == null) {
 synchronized (OldSingleton.class) {
 if(instance == null)
 instance = new OldSingleton();
 }
 }
 return instance;
 }
}
```

# Singleton pattern using enum type

- Create singleton using Java enum.

```
public enum MySingleton {
 INSTANCE;
 public void sayHello() {
 System.out.println("Hellod");
 }
 public void sayBye() {
 System.out.println("Bye");
 }
}
```

```
public class TestSingleton {
 public static void main(String[] args) {
 MySingleton singleton = MySingleton.INSTANCE;
 singleton.sayHello();
 singleton.sayBye();
 }
}
```

# Avoid creating unnecessary objects

- If this usage occurs in a loop or in a frequently invoked method, millions of `String` instances can be created needlessly

```
String s = new String("stringette");
```

- This code guarantees that the object will be reused

```
String s = "stringette";
```



# Avoid creating unnecessary objects

- The isBabyBoomer method unnecessarily creates a new Calendar, TimeZone, and two Date instances each time it is invoked

```
public class Person {
 private final Date birthDate = null;

 // DON'T DO THIS!
 public boolean isBabyBoomer() {
 // Unnecessary allocation of expensive object
 Calendar gmtCal = Calendar.getInstance(TimeZone.getTimeZone("GMT"));
 gmtCal.set(1946, Calendar.JANUARY, 1, 0, 0, 0);
 Date boomStart = gmtCal.getTime();
 gmtCal.set(1965, Calendar.JANUARY, 1, 0, 0, 0);
 Date boomEnd = gmtCal.getTime();
 return birthDate.compareTo(boomStart) >= 0
 && birthDate.compareTo(boomEnd) < 0;
 }
 public static void main(String[] args) {

 }
}
```

# Avoid creating unnecessary objects

- The version that follows avoids this inefficiency with a static initializer

```
Optimized_Person.java
package unnecessaryobjects;

import java.util.Date;

public class Optimized_Person {
 private final Date birthDate=null;
 private static final Date BOOM_START;
 private static final Date BOOM_END;
 static {
 Calendar gmtCal = Calendar.getInstance(TimeZone.getTimeZone("GMT"));
 gmtCal.set(1946, Calendar.JANUARY, 1, 0, 0, 0);
 BOOM_START = gmtCal.getTime();
 gmtCal.set(1965, Calendar.JANUARY, 1, 0, 0, 0);
 BOOM_END = gmtCal.getTime();
 }
 public boolean isBabyBoomer() {
 return birthDate.compareTo(BOOM_START) >= 0
 && birthDate.compareTo(BOOM_END) < 0;
 }
 public static void main(String[] args) {
 }
}
```

# Avoid creating unnecessary objects

- Change Long to long in code below:

```
public static void main(String[] args) {
 Long sum = 0L;
 for (long i = 0; i < Integer.MAX_VALUE; i++) {
 sum += i;
 }
 System.out.println(sum);
}
```

- Autoboxing is the automatic conversion that the Java compiler makes between the primitive types and their corresponding object wrapper classes. For example, converting an int to an Integer
- If the conversion goes the other way, this is called unboxing.

# Eliminate obsolete object references

- The program has a “memory leak,” which can silently manifest itself as reduced performance.

```
public class Stack {
 private Object[] elements;
 private int size = 0;
 private static final int DEFAULT_INITIAL_CAPACITY = 16;
 public Stack() {
 elements = new Object[DEFAULT_INITIAL_CAPACITY];
 }
 public void push(Object e) {
 ensureCapacity();
 elements[size++] = e;
 }
 public Object pop() {
 if (size == 0)
 throw new EmptyStackException();
 return elements[--size];
 }
 private void ensureCapacity() {
 if (elements.length == size)
 elements = Arrays.copyOf(elements, 2 * size + 1);
 }
}
```

# Eliminate obsolete object references

- NULL out references once they become obsolete

```
public class FixedStack {
 private Object[] elements;
 private int size = 0;
 private static final int DEFAULT_INITIAL_CAPACITY = 16;
 public FixedStack() {
 elements = new Object[DEFAULT_INITIAL_CAPACITY];
 }
 public void push(Object e) {
 ensureCapacity();
 elements[size++] = e;
 }
 public Object pop() {
 if (size == 0)
 throw new EmptyStackException();
 Object result = elements[--size];
 elements[size] = null; // Eliminate obsolete reference
 return result;
 }
 private void ensureCapacity() {
 if (elements.length == size)
 elements = Arrays.copyOf(elements, 2 * size + 1);
 }
}
```

# Using String literal instead of String object


- You can create String objects in the following ways.
- Way #1
  - `String s1 = "hello";`
  - `String s2 = "hello";`
- Way #2.
  - `String s3 = new String("hello");`
  - `String s4 = new String("hello");`
- Which of the above gives better performance?

# Using String literal instead of String object

- How the JVM works with Strings:
- Using == operator and String.equals() method.
  - == operator returns true if the references point to the same object
  - String.equals() method checks the contents of the String

# Using String literal instead of String object

- Optimization by Interning Strings
  - `String.intern()` method avoids duplicating String objects





# Optimization techniques when Concatenating Strings

- Compile time resolution versus run time resolution
  - Interestingly the + operator is faster than StringBuffer.append() method.
  - Before compilation:
    - `String result = "This is"+"testing the"+"difference"+"between"+"String"+"and"+"StringBuffer";`
  - After compilation
    - `String result = "This is testing the difference between String and StringBuffer";`
  - String object is resolved at compile time where as StringBuffer object is resolved at run time

# Optimization techniques when Concatenating Strings

## ■ Compile time resolution versus run time resolution

```
public class StringConcat {
 public static void main(String[] args) {
 // Test the String Concatination
 int count = 74836;
 long startTime = System.currentTimeMillis();
 for (int i = 0; i < count; i++) {
 String result = "This is" + "testing the" + "difference"
 + "between" + "String" + "and" + "StringBuffer";
 }
 long endTime = System.currentTimeMillis();
 System.out.println("Time taken for string concatenation using + operator : "
 + (endTime - startTime) + " milli seconds");
 // Test the StringBuffer Concatination
 long startTime1 = System.currentTimeMillis();
 for (int i = 0; i < count; i++) {
 StringBuffer result = new StringBuffer();
 result.append("This is");
 result.append("testing the");
 result.append("difference");
 result.append("between");
 result.append("String");
 result.append("and");
 result.append("StringBuffer");
 }
 long endTime1 = System.currentTimeMillis();
 System.out.println("Time taken for String concatenation using StringBuffer : "
 + (endTime1 - startTime1) + " milli seconds");
 }
}
```

- Time taken for string concatenation using + operator : 0 milli seconds
- Time taken for String concatenation using StringBuffer : 46 milli seconds

# Optimization techniques when Concatenating Strings

- Using StringBuffer instead of String

```
public class StringBufferTest {
 public static void main(String[] args) {
 // Test the String Concatenation using + operator
 long startTime = System.currentTimeMillis();
 String result = "hello";
 for (int i = 0; i < 1500; i++) {
 result += "hello";
 }
 long endTime = System.currentTimeMillis();
 System.out.println("Time taken for string concatenation using + operator : "
 + (endTime - startTime) + " milli seconds");

 // Test the String Concatenation using StringBuffer
 long startTime1 = System.currentTimeMillis();
 StringBuffer result1 = new StringBuffer("hello");
 for (int i = 0; i < 1500; i++) {
 result1.append("hello");
 }
 long endTime1 = System.currentTimeMillis();
 System.out.println("Time taken for string concatenation using StringBuffer : "
 + (endTime1 - startTime1) + " milli seconds");
 }
}
```

- Time taken for string concatenation using + operator : 16 milli seconds
- Time taken for string concatenation using StringBuffer : 0 milli seconds

# Class and Instance Initialization

```
class cls_init1 {
 static class Data {
 private int month;
 private String name;
 Data(int i, String s) {
 month = i;
 name = s;
 }
 }
 Data months[] = {
 new Data(1, "January"),
 new Data(2, "February"),
 new Data(3, "March"),
 new Data(4, "April"),
 new Data(5, "May"),
 new Data(6, "June")
 };
 public static void main(String args[]) {
 final int N = 250000;
 cls_init1 x;
 Timer t = new Timer();
 for (int i = 1; i <= N; i++)
 x = new cls_init1();
 t.print("data declared non-static");
 }
}
```

```
class Timer {
 long t;
 public Timer() {
 reset();
 }
 public void reset() {
 t = System.currentTimeMillis();
 }
 public long elapsed() {
 return System.currentTimeMillis() - t;
 }
 public void print(String s) {
 System.out.println(s + ": " + elapsed());
 }
}
```

# Class and Instance Initialization

- Turn the number/name data into a class variable, with a single copy across all instances

```
class cls_init2 {
 static class Data {
 private int month;
 private String name;
 Data(int i, String s) {
 month = i;
 name = s;
 }
 }
 static Data months[] = {
 new Data(1, "January"),
 new Data(2, "February"),
 new Data(3, "March"),
 new Data(4, "April"),
 new Data(5, "May"),
 new Data(6, "June")
 };
 public static void main(String args[]) {
 final int N = 250000;
 cls_init2 x;
 Timer t = new Timer();
 for (int i = 1; i <= N; i++)
 x = new cls_init2();
 t.print("data declared static");
 }
}
```

# Inner Classes

- JVM has restrictions on calling private members from outside of their class.
- A special access method is generated by the compiler and added internally to the meth\_inner class

```
public class meth_inner {
 private void f() {
 System.out.print("hello!");
 }
 class A {
 A() {
 f();
 }
 }
 public meth_inner() {
 A a = new A();
 }
 public static void main(String args[]) {
 meth_inner x = new meth_inner();
 }
}
```

# Inner Classes

- A generated method has a name `access$0()`, and it in turns calls `f()`.
- If you use the JDK utility program that disassembles `.class` files, by saying:

```
$ javap -c meth_inner
```

- The output includes the following method definition:

```
Method void access$0(meth_inner)
```

```
0 aload_0
```

```
1 invokespecial #7
```

```
4 return
```

# Using Number instead of Strings

- Creating a Double from a string takes about 15 times as long as from a number.

```
public class meth_inner {
 public static void main(String args[]) {
 final int N = 100000;
 Double d;
 Timer t = new Timer();
 for (int i = 1; i <= N; i++)
 d = new Double(12.34);
 t.print("as number");

 t.reset();
 for (int i = 1; i <= N; i++)
 d = new Double("12.34");
 t.print("as string");
 }
}
```



# Loops

- Most people write their for loops like this:

```
for (i=0; i<n; i++) {
 // do some stuff
}
```

- But, since in almost every language, comparing an int to 0 is faster.

```
for (i=n-1; i>=0; i--){
 // do some stuff
}
```

- You don't have to do the subtraction at the beginning of the loop.

```
for (i=n; --i>=0;) {
 // do some stuff
}
```

# Java Optimization Tip #1

- Data Structures
  - Have a profound influence on performance
    - Early design helps once
  - Choice of data structure can constrain what algorithms you can use
- Proportionality to Caller
  - Foo() takes 1 ms. Bar() calls foo.
    - If Bar() takes 20 ms, it's not worth looking at Foo()
    - If Bar() takes 2ms, then we should look at Foo()

# Java Optimization Tip #2

- 1-1 User Event Rule
  - If something happened a fixed number of times (1-3) for each user event, then it's not worth looking at
  - If something happens 100s of times for each user even then it is worth looking at
- 1-10-100 Rule
  - Assignment – 1 unit of time
  - Method call – 10 units of time
  - New Object or Array – 100 units of time
    - Rule of thumb only. Not scientific.
    - Hard to determine the actual cost
  - Bad idea to try and maintain your own free list. The GC knows best.

# Java Optimization Tip #3

- `int getWidth()` vs. `Dimension getSize()`
  - `getSize()` requires a heap allocated object
  - `getWidth()` and `getHeight()` may just be inlined to move the two ints right into the local variables of the caller code

# Java Optimization Tip #4

- Locals are faster than Instance variables
  - Local (stack) variables faster than any member variables of objects
    - Easier for the optimizer to work with
- Inside loops, pull needed values into local variables
  - 1. Slow: message send
    - ... `i < piece.getWidth()`
  - 2. Medium: instance variable
    - ... `i < piece.width`
  - 3. Fast: local variable
    - ... `final int width = piece.getWidth`
    - ... `i < width`
      - This is faster since the JIT can put the value in a native register

# Java Optimization Tip #4

- static (class) methods
  - These are the fastest to call, taking around 220ns.
- final methods
  - These are somewhere between static and instance methods, taking around 300ns.
- instance methods
  - These are a little slower, taking around 550ns.
- interface methods
  - These are surprisingly slow, taking on the order of 750ns to call.
- synchronized methods
  - These are by far the slowest, since an object lock has to be obtained, and take around 1,500ns.

# Java Optimization Tip #5

- Avoid Synchronized (Vector)
  - Synchronized methods have a cost associated with them
    - This is significantly improved in Java 1.3
  - Can have synchronized and unsynchronized methods and switch based on some flag
  - Use “immutable” objects to finesse synchronization problems
  - *Immutable* object cannot be changed after it is constructed.
  - Vector class is synchronized for everything
    - Use ArrayList instead!
    - If you can use a regular array, even better

# Java Optimization Tip #6

- Don't Parse
  - Obvious but slow strategy – read in XML, ASCII, etc.
  - Build a big data structure
- Faster approach
  - Read into memory, but keep as characters
  - Search/Parse when needed
  - Or Parse only subparts



# Java Optimization Tip #7

- Avoid weird code
  - JVM will optimize most standard coding styles
    - So write code in the most obvious, common way
  - Weird code is often the result of an attempt at optimization!
- Let the JIT/Hotspot do its thing!

# Java Optimization Tip #8

- Threading / GUI Threading
  - Use separate thread to ensure the GUI is *snappy*
- Pros
  - Makes best use of parallel hardware
- Cons
  - Software is harder to write
  - Bugs can be subtle
  - Locking/Unlocking costs

# Java Optimization Tip #9

- A local variable assigned only once can be declared final.
- A method argument that is never assigned can be declared final.
- Detects when a new object is created inside a loop
- ArrayList is a much better Collection implementation than Vector.
- AddEmptyString

*String s = "" + 123; // bad*

*String t = Integer.toString(456); // ok*

# References

- Effective Java
  - Joshua Bloch
- Performance improvement techniques in String and StringBuffer
  - [comments@precisejava.com](mailto:comments@precisejava.com)
- Java™ Performance Tuning and Java Optimization Tips
  - [Glen McCluskey](#)